

Suggest up to three different algorithms for computing the join of two tables T1 and T2. Contrast the algorithms in terms of their time complexity and storage requirements.

POSSIBLE ANSWERS:

Algorithm 1: simple nested-loop join

```
for each row R1 in table T1:
  for each row R2 in table T2:
    if R1 and R2 satisfy the joining condition then combine
    R1 and R2 and append to the result table
```

Time complexity: $O(|T1| * |T2|)$

Storage requirement: essentially none

Algorithm 2: single-sort join

```
sort table T1 on the joining value;
for each row R2 in T2:
  use binary search on sorted T1 to find all matching values and
  add to the result table
```

Time complexity: $O(|T1| * \log(|T1|))$ to sort T1
 $O(|T2| * \log(|T1|))$ for second step

Algorithm 3: sort-merge join *Sort both*

```
sort table T1 on the joining value;
sort table T2 on the joining value;
traverse the two tables linearly (with some "backtracking" for
duplicate values), matching join values and adding joining tuples
to the result table
```

Time complexity: $O(|T1| * \log(|T1|))$ to sort T1
 $O(|T2| * \log(|T2|))$ to sort T2
 $O(\max(|T1|, |T2|))$ for "merge" phase assuming not too
many duplicate joining values

Storage requirement: not much, depends on sorting algorithm used



Algorithm : hash join

```
set up hash table;
for each row R1 in T1:
  hash R1's join value and put R1 in the appropriate hash bucket
for each row R2 in T2:
  hash R2's join value;
  find all matching tuples in the hash bucket and add to the result table
```

7

Time complexity: $|T1| + |T2|$ assuming well-distributed hash table

Storage requirement: $O(|T1|)$ for hash table