

Question for students with database implementation background

Consider the standard tuple-based nested-loop join algorithm for computing  $T_1 \text{ JOIN } T_2$ :

```
for each row R1 in T1:
  for each row R2 in T2:
    if R1,R2 satisfy the join condition then add R1/R2 to result
```

Suggest three separate possible improvements to this algorithm. Assume a standard DBMS storage system and query processing context. Improvements could depend on additional assumptions or scenarios.

ANSWER:

1. Nested-block join

Process  $T_1$  and  $T_2$  block-at-a-time instead of row-at-a-time. Considerably reduces the number of times  $T_2$  is scanned, without incurring extra I/O's for  $T_1$ . ★

"block at a time"

2. "Rocking"

For  $T_2$ , scan it forwards the first time, backwards the second time, forwards the third time, etc. Takes advantage of LRU page replacement policy typically used by database buffer managers.

forward - backward

last recent use

3. Use of keys

If the join condition is  $T_1.A = T_2.B$  and  $B$  is a key for  $T_2$ , then once a match is found the algorithm can break out of the inner loop.

4. Use of index

If the join condition is  $T_1.A = T_2.B$  and there is an index on  $T_2.B$  then the inner loop can find matching  $T_2$  rows using the index instead of by scanning the whole relation. (Also works for inequality join conditions if the index is a B-tree.)